

Very brief introduction to Lucene (4.10)



[From yonik@apache.org, 02 May 2007]

Amsterdam, Netherlands, slides:
<http://www.apache.org/~yonik>

Related Projects

[Apache Hadoop](#)

[Apache ManifoldCF](#)

[Apache Lucene.Net](#)

[Apache Lucy](#)

[Apache Mahout](#)

[Apache Nutch](#)

[Apache OpenNLP](#)

[Apache Tika](#)

[Apache Zookeeper](#)

What is Lucene

- High performance, scalable, full-text search library
- Focus: Indexing + Searching Documents
- 100% Java, no dependencies, no config files
- **No crawlers or document parsing**
- Users: Wikipedia, Technorati, Monster.com, Nabble, TheServerSide, Akamai, SourceForge, Twitter, LinkedIn, Hi5, ...
- Applications: Eclipse, JIRA, Roller, OpenGrok, Nutch, Solr, and many commercial products.

Main features

Scalable, High-Performance Indexing

over 150GB/hour on modern hardware

small RAM requirements -- only 1MB heap

incremental indexing as fast as batch indexing

index size roughly 20-30% the size of text indexed

Cross-Platform Solution

Available as Open Source software under the

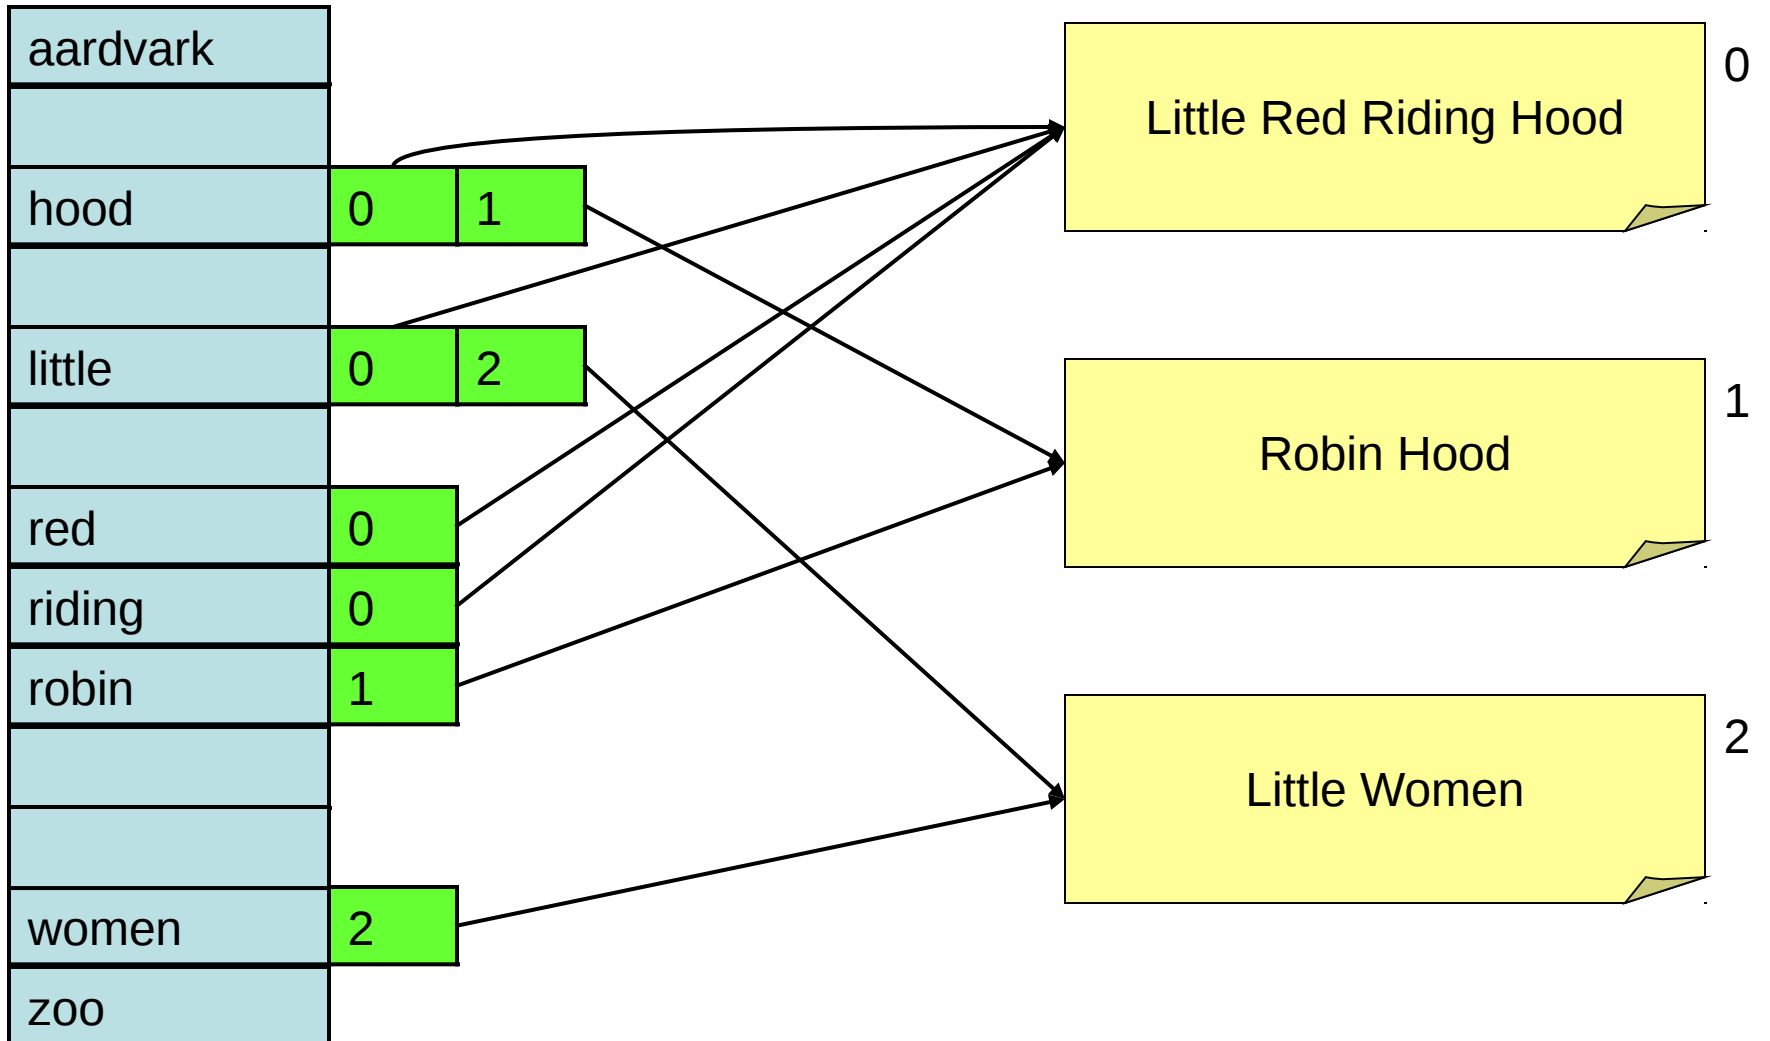
Apache License which lets you use Lucene in both commercial and Open Source programs

100%-pure Java

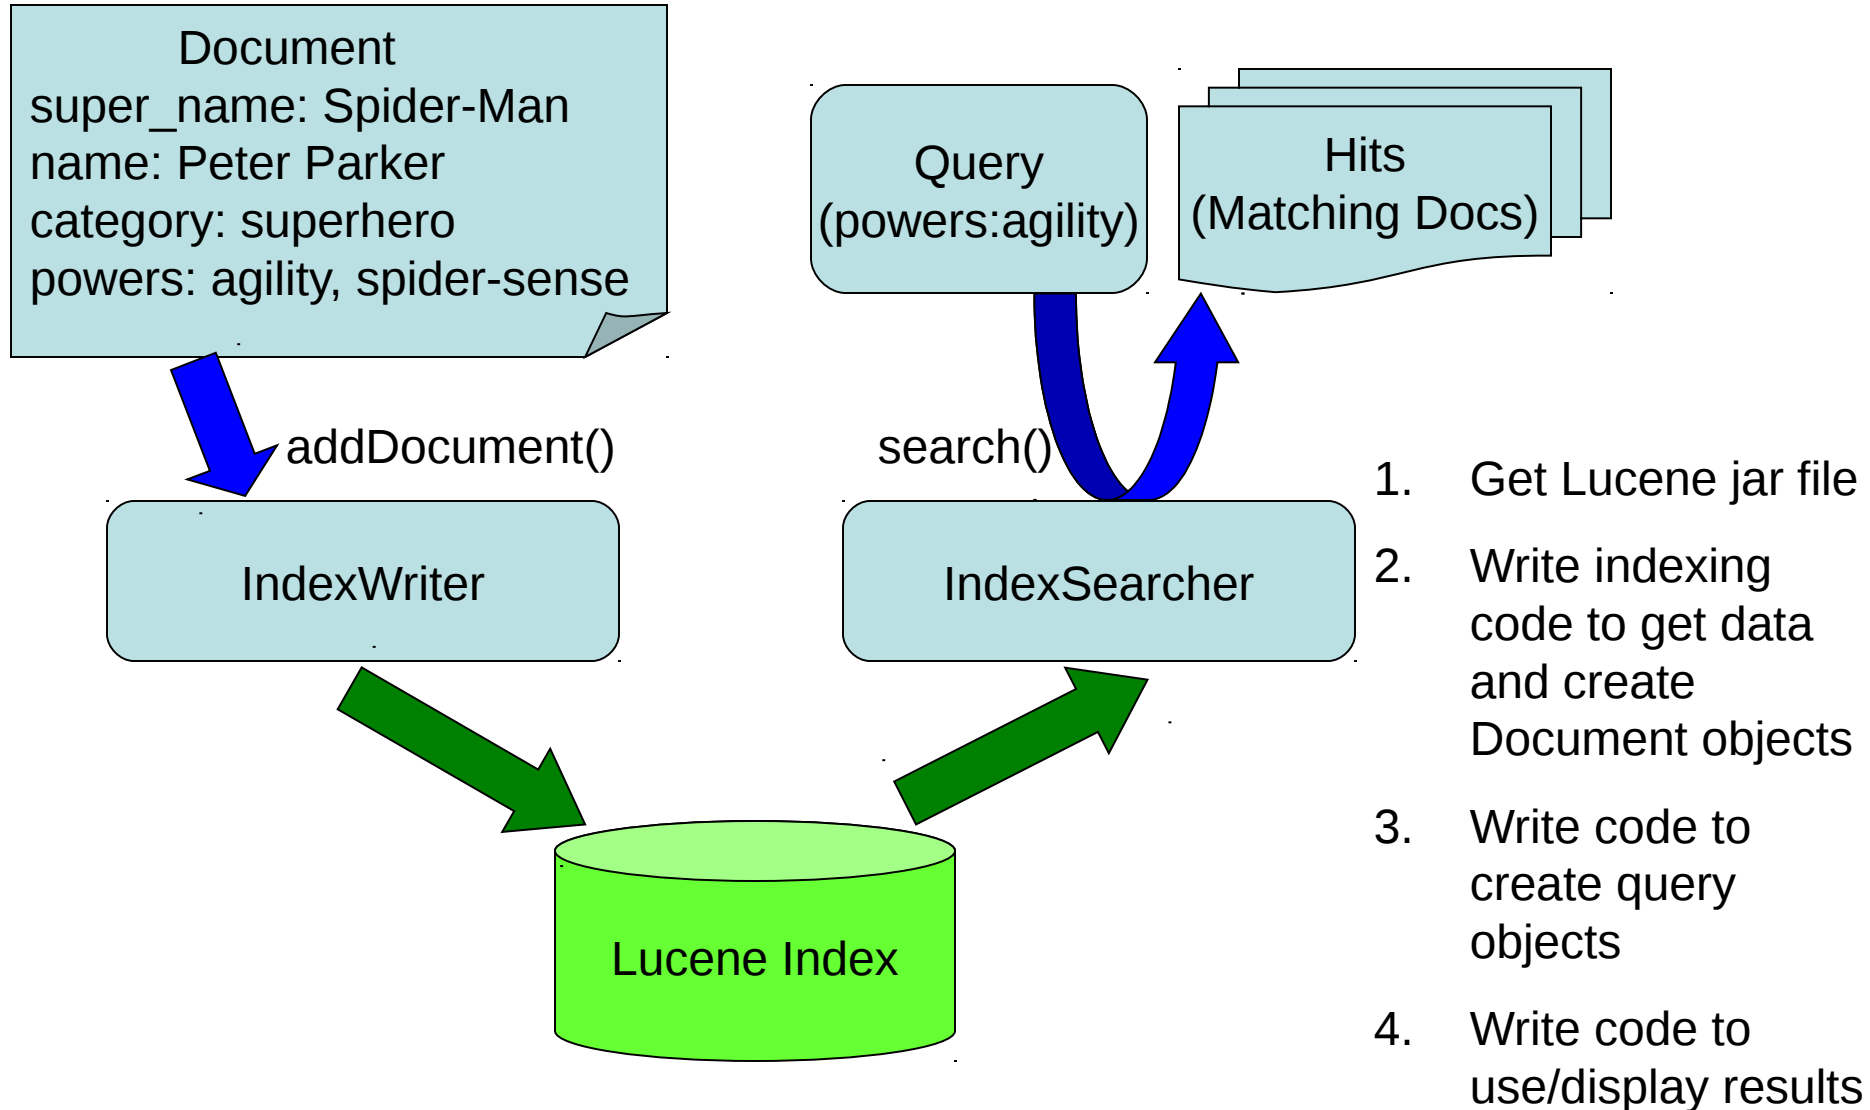
Implementations

in other programming languages available that are index-compatible

Inverted Index



Basic Application



Indexing Documents

```
IndexWriter writer = new IndexWriter(directory, analyzer,  
    true);  
Document doc = new Document();  
doc.add(new Field("super_name", "Sandman",  
    Field.Store.YES, Field.Index.TOKENIZED));  
doc.add(new Field("name", "William Baker",  
    Field.Store.YES, Field.Index.TOKENIZED));  
doc.add(new Field("name", "Flint Marko",  
    Field.Store.YES, Field.Index.TOKENIZED));  
// [...]  
writer.addDocument(doc);  
writer.close();
```

Field Options

- **Indexed**
 - Necessary for searching or sorting
- **Tokenized**
 - Text analysis done before indexing
- **Stored**
 - You get these back on a search “hit”
- **Compressed**
- **Binary**
 - Currently for stored-only fields

Searching an Index

```
IndexSearcher searcher = new IndexSearcher(directory);
QueryParser parser = new QueryParser("defaultField",
    analyzer);
Query query = parser.parse("powers:agility");
Hits hits = searcher.search(query);
System.out.println("matches:" + hits.length());
Document doc = hits.doc(0); // look at first match
System.out.println("name=" + doc.get("name"));
searcher.close();
```


Scoring

- **VSM** – Vector Space Model
- **tf** – term frequency: number of matching terms in field
- **lengthNorm** – number of tokens in field
- **idf** – inverse document frequency
- **coord** – coordination factor, number of matching terms

Score Boosting

Lucene allows influencing search results by "boosting" in more than one level:

- Document level boosting - while indexing - by calling **document.setBoost()** before a document is added to the index.
- Document's Field level boosting - while indexing - by calling **field.setBoost()** before adding a field to the document (and before adding the document to the index).
- Query level boosting - during search, by setting a boost on a query clause, calling **Query.setBoost()**.

Lucene Conceptual Scoring Formula

$\text{score}(q,d) =$

$\text{coord-factor}(q,d) \cdot \text{query-boost}(q) \cdot$

$(V(q) \cdot V(d) / |V(q)|) \cdot \text{doc-len-norm}(d) \cdot \text{doc-boost}(d)$

Query Construction

Lucene QueryParser

- Example: *queryParser.parse("name:Spider-Man");*
- Easy entered queries, debugging, IPC
- Does text analysis and constructs appropriate queries
- Not all query types supported

Programmatic query construction

- Example:
new TermQuery(new Term("name","Spider-Man"))
- Explicit
- Does not do text analysis for you

Query Examples

1. justice league

- EQUIV: justice OR league
- QueryParser default is “optional”

2. +justice +league –name:aquaman

- EQUIV: justice AND league NOT name:aquaman

3. “justice league” –name:aquaman

4. title:spiderman^10 description:spiderman

5. description:“spiderman movie”~10

Query Examples 2

1. **releaseDate:[2000 TO 2007]**

- Range search: lexicographic ordering, so beware of numbers

2. **Wildcard searches:** sup?r, su*r, super*

3. **spider~**

- Fuzzy search: Levenshtein distance
- Optional minimum similarity: spider~0.7

4. **(Superman AND “Lex Luthor”) OR (+Batman +Joker)**

Deleting Documents

Deleting with IndexWriter

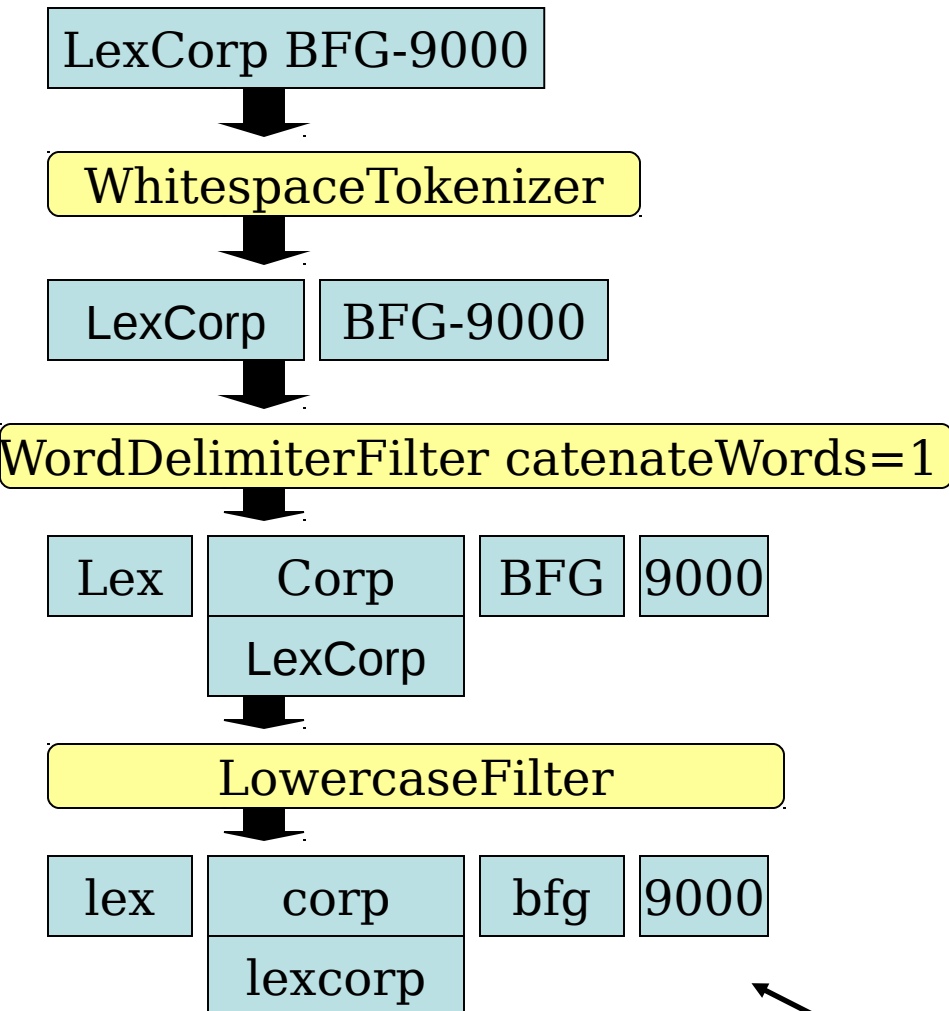
- **deleteDocuments**(Term term): *Deletes the document(s) containing term.*
- **updateDocument**(Term term, Iterable<? extends IndexableField> doc): *Updates a document by first deleting the document(s) containing term and then adding the new document.*
- Deleting does not immediately reclaim space

Performance

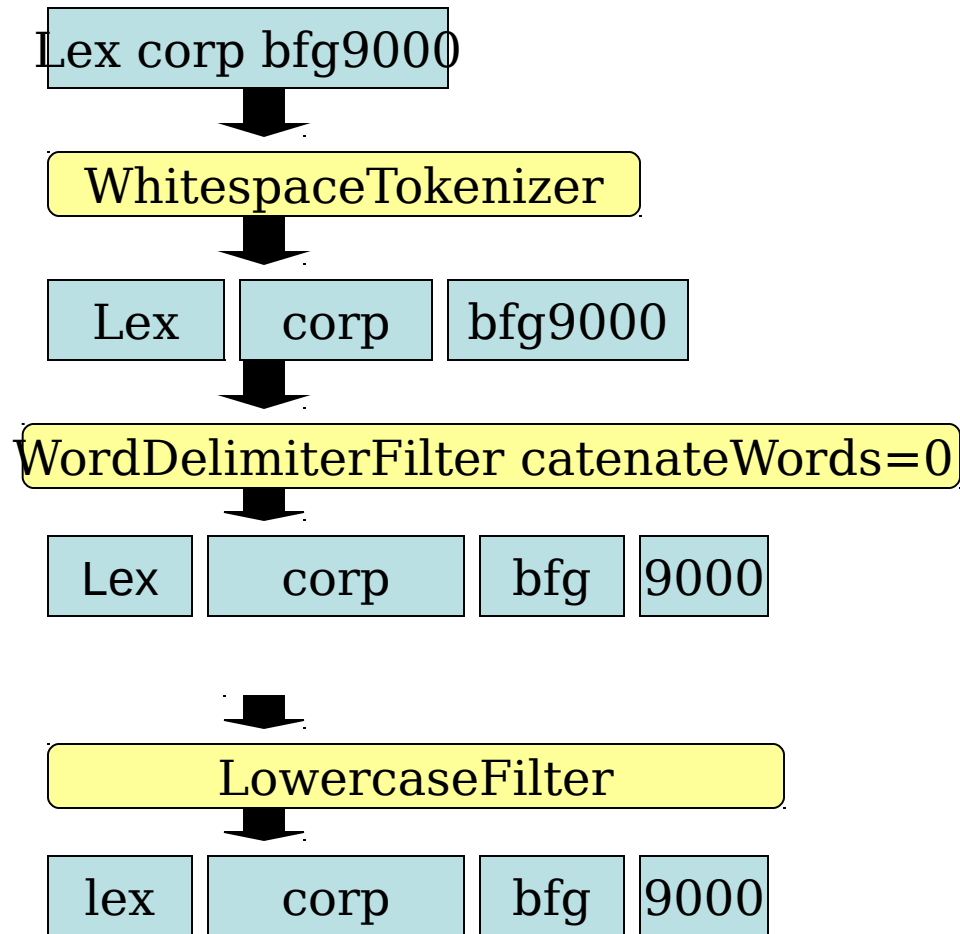
- **Indexing Performance**
 - Index documents in batches
 - Raise merge factor
 - Raise maxBufferedDocs
- **Searching Performance**
 - Reuse IndexSearcher
 - Optimize
 - Use cached filters (see QueryFilter)
 - ‘+superhero +lang:english’
 - ‘superhero’ filtered by ‘lang:english’

Analysis & Search Relevancy

Document Indexing Analysis



Query Analysis



A Match!

Tokenizers

Tokenizers break field text into tokens

- **StandardTokenizer**
 - source string: “full-text lucene.apache.org”
 - “full” “text” “lucene.apache.org”
- **WhitespaceTokenizer**
 - “full-text” “lucene.apache.org”
- **LetterTokenizer**
 - “full” “text” “lucene” “apache” “org”

TokenFilters

- LowerCaseFilter
- StopFilter
- ISOLatin1AccentFilter
- SnowballFilter
 - stemming: reducing words to root form
 - rides, ride, riding => ride
 - country, countries => countri
- contrib/analyzers for other languages
- SynonymFilter (from Solr)
- WordDelimiterFilter (from Solr)

Analizers

```
class MyAnalyzer extends Analyzer {  
    private Set myStopSet =  
        StopFilter.makeStopSet(StopAnalyzer.ENGLISH_STOP_WORDS);  
  
    public TokenStream tokenStream(String fieldname, Reader reader) {  
        TokenStream ts = new StandardTokenizer(reader);  
        ts = new StandardFilter(ts);  
        ts = new LowerCaseFilter(ts);  
        ts = new StopFilter(ts, myStopSet);  
        return ts;  
    }  
}
```

Analysis Tips

- Use *PerFieldAnalyzerWrapper*

This analyzer is used to facilitate scenarios where different fields require different analysis techniques.

- Add same field more than once, analyze differently
 - Boost exact case matches
 - Boost exact tense matches
 - Query with or without synonyms
 - Soundex for sounds-like queries
- Use *explain(Query q, int docid)* for debugging

Nutch

- Open source web search application
- Crawlers
- Link-graph database
- Document parsers (HTML, word, pdf, etc)
- Language + charset detection
- Utilizes Hadoop (DFS + MapReduce) for massive scalability

Solr

Solr™ is the popular, blazing fast open source enterprise search platform from the Apache Lucene™ project.

- REST XML/HTTP, JSON APIs
- Faceted search
- Flexible Data Schema
- Hit Highlighting
- Configurable Advanced Caching
- Replication
- Web admin interface

Running Toy Application

- <http://www.lucene-tutorial.com/code/TextFileIndexer.java>
- LUCENE 4.5 <http://mirror.sdunix.com/apache/lucene/java/4.5.0/>
- Importing external JARs
 - ../../lucene-4.5.0/core/lucene-core-4.5.0.jar
 - ../../lucene-4.5.0/analysis/common/lucene-analyzers-common-4.5.0.jar
 - ../../lucene-4.5.0/queryparser/lucene-queryparser-4.5.0.jar